

Support Vector Machine (SVM)

Image classification

- Image classification is the task of ***assigning a label to an image*** from a ***pre-defined set of categories***.
 - Example:
labels {cat, dog, bear, fish, bird}



Label: dog

Image classification is a challenging problem

- We see “aleph”



Computer sees a matrix of numbers

177	180	175	169	179	179	166	111	160	180
140	148	165	163	179	180	153	63	105	179
84	57	104	154	170	174	154	64	84	174
140	96	60	72	87	101	141	77	83	173
182	172	97	42	49	54	74	71	72	164
181	156	63	52	111	142	103	58	61	156
185	140	52	115	171	177	171	113	55	137
179	129	49	132	180	176	168	154	83	132
166	132	53	130	182	179	169	166	112	92
173	170	129	154	183	177	169	168	141	125

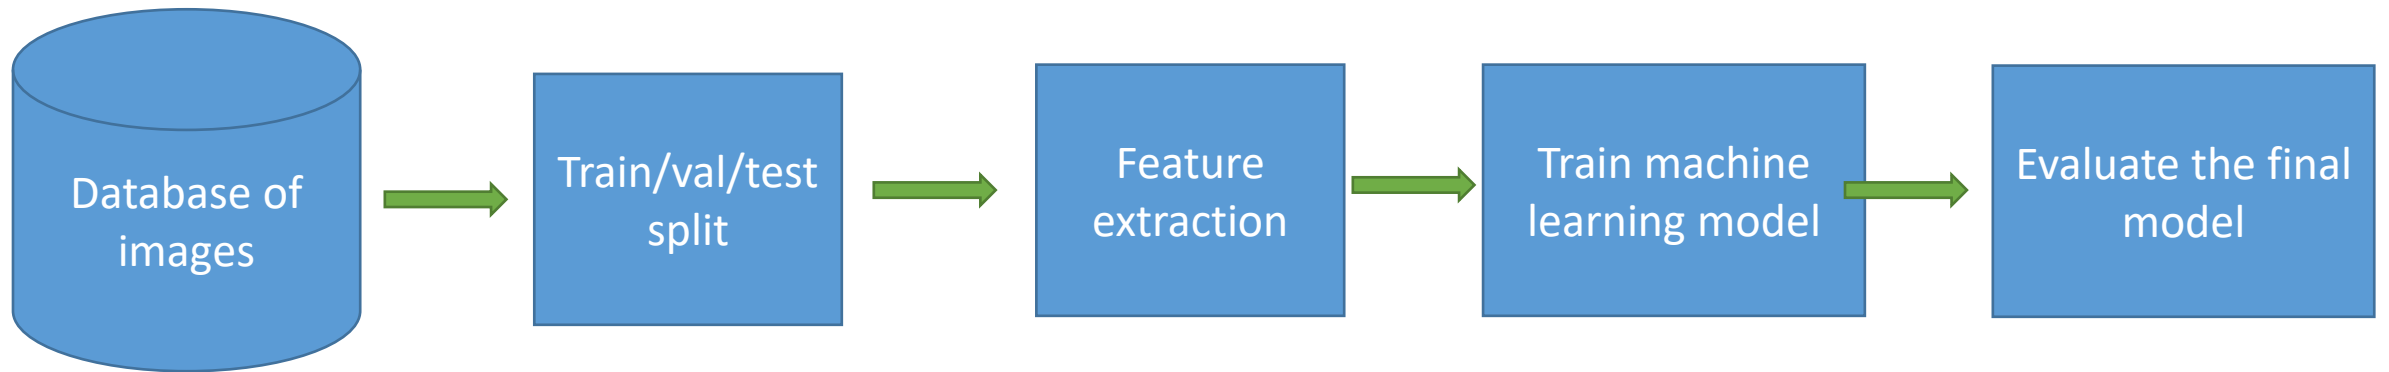
Machine learning for image classification

- Describe an image(image part) by a descriptor
- Apply machine learning to “teach” the computer how objects from each class look
 - Supervised learning
 - Unsupervised learning
 - Semi-supervised learning

Machine learning

- How to tell the difference between a dog and other categories?
- Instead of a coding set of rules to recognize each class, in machine learning we use a **data-driven approach**
 - Supply examples of each class and teach the model to recognize between them
 - Training set: each example consist of a pair (image, label)

The image classification pipeline



Compiling the dataset

- Collect the images from each class
- The dataset should be balanced – approximately the same number of images for each class

Train/Val/Test Sets

- Training set – examples used to train the model ($\approx 80\%$)
- Validation set – a “test set” used in training to tune various hyperparameters ($\approx 10\%$)
- Test set – evaluate the performance of the model ($\approx 10\%$)

Feature extraction



→ [4, 0.25, 1.7,]



→ [3, 0.37, 7.1,]

Training

- Machine learning models
 - Support Vector Machine (SVM)
 - Decision Trees
 - Random Forest
 - K-Nearest Neighbor
 - Neural networks
 - ...

Evaluation


- For each of the images in the test set we predict the label and compare to its ground-truth (actual) label
- Calculate various metrics based on the results, e.g., precision, recall, f-measure

Support Vector Machine (SVM)

- **Support-vector machines (SVMs)**, are supervised learning models with associated learning algorithms that analyze data for classification
- Developed at [AT&T Bell Laboratories](#) by [Vladimir Vapnik](#) with colleagues (1992)
- Given a set of training examples, each marked as belonging to one of two categories, an SVM training algorithm builds a model that assigns new examples to one category or the other

Support Vector Machine (SVM)


- SVM maps training examples to points in space so as to maximize the width of the gap between the two categories
- New examples are then mapped into that same space and predicted to belong to a category based on which side of the gap they fall




Refence [wikipedia](https://en.wikipedia.org/wiki/Support_vector_machine)

Separating Hyperplane




- **Primary assumption:** feature vectors of similar images lie *close together* in an n -dimensional space
- The line used to separate the data is called the separating hyperplane.




Separating Hyperplane



Separating Hyperplane



Maximum-Margin Separating Hyperplane



Maximum-Margin Separating Hyperplane - largest distance to the nearest training-data point of any class

Non-linearly separable datasets



We cannot separate these two sets by a straight line

The kernel trick

- Project all the points into a higher dimension using so-called a **kernel matrix**




Image taken from <https://towardsdatascience.com/the-kernel-trick-c98cdbcaeb3f>

Types of kernel

- Linear: $K(x, y) = x^T y$
- Polynomial: $K(x, y) = (x^T y + c)^d$
- Radial basis function kernel: $K(x, y) = \exp(-\gamma ||x - y||^2)$
- Sigmoid: $K(x, y) = \tanh(\gamma(x^T y) + r)$

x and y are vectors in the *input space*, i.e. vectors of features computed from training or test sets

SVM implementation

- [Scikit-learn](#) library

```
from sklearn import svm
```

```
from sklearn.metrics import classification_report
```

```
# Create a linear SVM classifier  
clf = svm.SVC(kernel='linear')
```

```
# Train classifier  
clf.fit(train_data, train_labels)
```

```
# Make predictions on unseen test data  
clf_predictions = clf.predict(test_data)  
print("Accuracy: {}".format(clf.score(test_data, test_labels) * 100))
```

```
# printing report and different statistics  
print(classification_report(test_labels, clf.predict(test_data)))
```

SVM implementation

- [Scikit-learn](#) library

```
from sklearn import svm
```

```
from sklearn.metrics import classification_report
```

```
# Create a linear SVM classifier  
clf = svm.SVC(kernel='linear')
```

```
# Train classifier  
clf.fit(train_data, train_labels)
```

```
# Make predictions on unseen test data  
clf_predictions = clf.predict(test_data)  
print("Accuracy: {}".format(clf.score(test_data, test_labels) * 100))
```

```
# printing report and different statistics  
print(classification_report(test_labels, clf.predict(test_data)))
```